

LDSSA

SLU16 Workflow

December 3rd, 2023



1. Introduction

Motivation

- **Data Science** is largely an **engineering discipline**. Writing code is an engineering practice and data science is mostly done with code.
- **A healthy respect for the engineering element of data science is key**. It will make your life an order of magnitude easier, and more importantly, it makes your **data science more responsible**.



Overview

In this **SLU** we will be covering the following:


- Workflow
- Pipelines and Custom Objects
- Workflow tips tricks
- Best practices





2. Topic Explanation


| Workflow |



Workflow

Specialization #1 - SLU04 & SLU05

1. Get the Data
2. Data Analysis and Preparation
 - a. Data analysis
 - b. Dealing with data problems
 - c. Feature engineering
 - d. Feature selection
3. Train Model
4. Evaluate Results
5. Iterate



| Pipelines and Custom Objects |



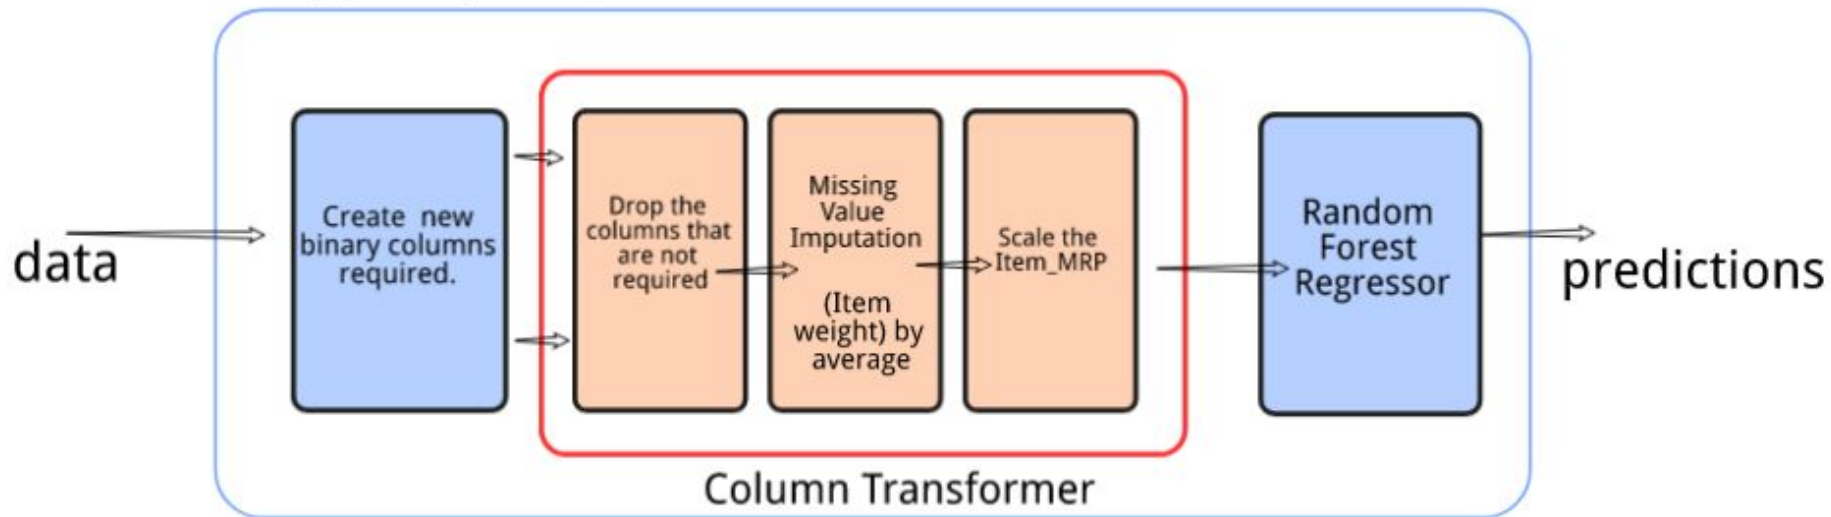
Sklearn Pipeline

- “Sequentially apply a list of transforms and a final estimator. “
- It implements the same API as the models (has ``predict`` and/or ``predict_proba``) but it applies each of the steps before calling the model with the input!



pipeline.fit() when using the training data

pipeline.predict() on test data



Pipeline

sklearn.pipeline

```
from sklearn.pipeline import Pipeline
from sklearn.feature_extraction.text import (
    CountVectorizer, TfidfTransformer)
from xgboost import XGBClassifier

stopwords, lemmatizer = ...

pipeline = Pipeline([
    ('preprocess', MessagePreprocessor(subject_weight=2)),
    ('text', TextProcessor(stopwords, lemmatizer)),
    ('vect', CountVectorizer()),
    ('tfidf', TfidfTransformer()),
    ('clf', XGBClassifier(objective='multi:softmax')),
])
```



Different objects

The main objects in scikit-learn are (one class can implement multiple interfaces):



Estimator:

The base object, implements a `fit` method to learn from data, either:

```
estimator = estimator.fit(data, targets)
```

or:

```
estimator = estimator.fit(data)
```

Predictor:

For supervised learning, or some unsupervised problems, implements:

```
prediction = predictor.predict(data)
```

Classification algorithms usually also offer a way to quantify certainty of a prediction, either using `decision_function` or `predict_proba`:

```
probability = predictor.predict_proba(data)
```



Transformer:

For filtering or modifying the data, in a supervised or unsupervised way, implements:

```
new_data = transformer.transform(data)
```

When fitting and transforming can be performed much more efficiently together than separately, implements:

```
new_data = transformer.fit_transform(data)
```




Model:

A model that can give a `goodness of fit` measure or a likelihood of unseen data, implements (higher is better):

```
score = model.score(data)
```

| Workflow tips tricks |





1. Establish a simple baseline FAST
2. Incrementally increase complexity
3. Use (and abuse) Scikit pipelines & Custom transformers

| Best Practises |





- **Use well-named variables**
- **Use functions!**
 - If you take a parameter that is a dataframe, always name it `_df`
 - Outside of functions, never name a dataframe `_df`
 - Inside of functions, never name a dataframe `df`
- **Immutability is key!**
 - Never use `inplace=True`
 - Always copy a dataframe at the beginning of the function and make your changes to that copy.



- **Imports up top!**
- **Organize your directory**
- **When writing functions**
 - Keep them close to the cell where you are testing it at first
 - When they are stabilizing, move them to the top of the notebook
 - When you are starting to use them a lot, move them into a `utils.py` file and import from there
- **“Restart and Run All”**
- **Each experiment should be runnable in a single cell**



3. Recap

- “First do it, then do it well, then do it better”
- Be organized - keep track of the improvements you need/want to do
- Try to use Scikit Pipelines whenever you can
- Make sure your code is readable not just to you, but to others!



```
8 // Dear programmer:
9 // When I wrote this code, only god and
10 // I knew how it worked.
11 // Now, only god knows it!
12 //
13 // Therefore, if you are trying to optimize
14 // this routine and it fails (most surely),
15 // please increase this counter as a
16 // warning for the next person:
17 //
18 // total_hours_wasted_here = 254
19 //
20
```



4. Q&A